

Graph-based models of volumetric medical data and applications to topological correction

Quentin Fortier

August 10, 2010

Contents

1	Introduction	2
1.1	Overview	2
1.2	Modelling	2
1.3	NP-hardness	4
1.4	Relations between the two problems	4
2	Previous works on surface cutting	4
2.1	Cutting a surface to reduce its genus to 0	5
2.2	Shortest homotopic cycles on a surface	5
2.3	First attempts to a volume cutting	6
3	General algorithm	7
3.1	Overview	7
3.2	First decomposition	7
3.3	Pant(s) optimization	8
3.3.1	Optimization of an auto 2-cut \mathcal{C}	8
3.3.2	Optimization of a 2-cut \mathcal{C} (not auto)	9
3.4	Processing choice	10
3.5	Selection of a cutting system	10
4	Improvements	10
4.1	Neighbourhood	10
4.2	Processing order and intersections	11
4.3	Complexity	11
5	Implementation	11
5.1	Overview	11
5.2	SGL	12
5.3	Data structures	12
5.4	Core algorithm	13
5.5	Orientation	13
5.6	Python script	14
5.7	Running time comparison	14

6	Appendix	15
6.1	Mathematical background	15
6.2	Applications	16
6.3	SGL	16
Todo: Ex ENS, diff tetgen -q, precalculer BFS pour optimize all		

1 Introduction

The reader not familiar with basic algebraic topology will find in the appendix the definitions needed in this report.

1.1 Overview

The limited resolution or the noise may increase the genus of the surfaces given by medical images and we want to remove them, both to have a more accurate image and apply algorithms that need genus 0 surfaces.

For that purpose, important works have been done in finding the shortest system of loops cutting reducing the genus to 0: several methods were proposed, for example in [1] or [2].

The goal of this internship is to extend these methods to a volumic topological correction, in which only few works have been done.

1.2 Modelling

For more precision, we first need some definitions: (all surfaces are assumed to be closed)

Definition 1. A *volume* \mathbb{V} (resp. *surface*) is a topological 3(resp. 2)-manifold. We assume a volume has one boundary $\partial\mathbb{V}$ and the genus g of \mathbb{V} (i.e the genus of $\partial\mathbb{V}$) is assumed to be different from 0.

From an algorithmic point of view, instead of an abstract 3-manifold we deal with either a 3-mesh (a set of tetrahedra) or voxels, which are obviously a volume in the sense of the above definition. We will refer to them as discrete volume. Discrete surfaces are either sets of triangles or of faces of a voxel set. Moreover, we assume these discrete volumes or surfaces to have a weight on their faces or edges (typically, the weight will be the length for an edge and the area for a face or the probability that the two volumic elements are adjacent, given by the segmentation step).




Figure 1: Voxels

Definition 2. Given a discrete volume \mathbb{V} , the *dual graph* $G^* = (V, E)$ of \mathbb{V} is constructed the following way:

- The vertices of G^* are the volumic elements (tetrahedra or voxels).
- Two vertices are connected if the corresponding two elementary volumes share a face, and the weight of this edge is the weight of the face.

Two edges of G^* are adjacent if the corresponding faces share an edge of \mathbb{V} .

Since every edge in G^* has a degree at most d , with $d = 6$ if \mathbb{V} is a voxel set and $d = 4$ if \mathbb{V} is a 3-mesh, $2|E| = \sum \deg(v) \leq d|V|$, so $O(E) = O(V)$.

Definition 3. Let A, B be two sets such that $A \subset B$ and B is connected. A is *B-non-separating* if $B - A$ is connected. Similarly, if $A_i \subset B$, $\{A_1, \dots, A_n\}$ is *B-non-separating* if $B - \bigcup A_i$ is connected.

Definition 4. A *1-cut* γ (on a surface \mathbb{S}) is a \mathbb{S} -non-separating loop.

On G^* , a 1-cut is a simple cycle of edges.

Remark 1. A contractible loop is always separating, but the converse is false. [cylinder]

Definition 5. A *2-cut* \mathcal{C} (in \mathbb{V}) is a connected, \mathbb{V} -non-separating, genus-0 surface such that $\partial\mathcal{C}$ lies on $\partial\mathbb{V}$.

On G^* , a 2-cut is a set of adjacent edges (however it is not sufficient, since we require the set to be a manifold).

Definition 6. The *cutting product* of a surface along a 1-cut γ is, informally, the surface produced by duplicating every edge and vertex and removing the adjacences between faces sharing an edge in γ .

The informal definition is similar for a volume.

See [1], p. 35, for a more precise definition.

Definition 7. A *system of 2 (resp. 1)-cuts* of \mathbb{V} (resp. \mathbb{S}) is a \mathbb{V} (resp. \mathbb{S})-non-separating set of 2 (resp. 1)-cuts such that the resulting cutting product has genus 0 and every pair of cuts doesn't intersect.

We can now give the main issue of this report:

Problem 1 (2-MIN-CUTTING). *Given a discrete volume \mathbb{V} , find a minimum system of 2-cuts.*

We are also interested in the related problem on surface, which we will investigate first:

Problem 2 (1-MIN-CUTTING). *Given a discrete surface \mathbb{S} , find a minimum system of 1-cuts.*

These minimums exist: there are finitely many sets of g 1-cuts (resp. 2-cuts).

1.3 NP-hardness

An apparently similar problem is known to be NP-hard:

Problem 3 (1-MIN-CUTTING-BOUNDARY). *Given a discrete surface \mathbb{S} , possibly with boundary, find a minimum cutting to obtain a single topological disk (i.e the resulting cutting product must have genus 0 and with no boundary).*

Theorem 1. *1-MIN-CUTTING-BOUNDARY is NP-hard*

Proof:

[7] reduces this problem to the *rectilinear Steiner tree problem* (i.e, given n points in the plane, find a tree connecting them all with only vertical and horizontal line segments), which is known to be NP-hard.

□

However, it is not known if 1-MIN-CUTTING and 2-MIN-CUTTING are NP-hard, but the two problems are linked.

In the case of voxels, we have:

Theorem 2. *1-MIN-CUTTING reduces to 2-MIN-CUTTING.*

Proof:

Let \mathbb{S} be an instance of 1-MIN-CUTTING: according to Jordan's theorem, \mathbb{S} bounds only one volume \mathbb{V} .

For each face of \mathbb{V} , we put a weight equals to sum of weights of the edges of this face that are in \mathbb{S} . This way the area of a surface in \mathbb{V} equals its perimeter.

The size of \mathbb{V} is polynomial in the size of \mathbb{S} according to the isoperimetric inequality: $36 \pi \text{vol}(\mathbb{V})^2 \leq \text{area}(\mathbb{S})^3$, and since the volume of \mathbb{V} equals the number of voxels in \mathbb{V} , the area of \mathbb{S} equals the number of faces (squares) in \mathbb{S} .

□

The same theorem holds for tetrahedra if we restrict the problem to PLC (*piece-wise linear complex*) which can be tetrahedrized (not all surfaces can be tetrahedrized).

1.4 Relations between the two problems

Although they have similarities, the two problems above are not identicals and especially, minimizing the area of a surface is not equivalent to minimizing its perimeter.

Moreover, a surface with bounded area may have a perimeter equals to infinity.

2 Previous works on surface cutting

In all this section, let \mathbb{S} be a surface with genus g .

This section introduce two algorithms dealing with **surface cutting**, the first described in [1], chap. 3 and the second in [2], chap. 3.

Both use a similar approach to find 1-cuts: they first show that computing a 1-cut is easy if we add constraints to it and then they split the surface, to force the wanted constraints.



Figure 2: Left: The Koch snowflake, a surface with finite area but with infinite perimeter. Right: The optimal 1-cut (in blue) is different from the optimal 2-cut (in red) of this torus.

2.1 Cutting a surface to reduce its genus to 0

Using the following approach, we can get a good approximation of a minimum system of 1-cuts of \mathbb{S} .

Lemma 1. We can calculate the shortest 1-cut on \mathbb{S} passing through a fixed point p in $O(V \log(V))$.

Proof:

We use Dijkstra's algorithm from p . As soon as we meet a vertex previously visited, we look at the number c of connected components of the complementary of the surface: if c equals one we continue our search, otherwise (c equals two) we found a shortest 1-cut from p .

□

Lemma 2. We can compute a first system of 1-cut \mathcal{Y} on \mathbb{S} in $O(V)$.

Proof:

Let t be a volumic element (triangle). We use the *cut locus* associated to t , i.e. the set of all points having more than one shortest path from t .

□

Lemma 3. Every 1-cut of \mathbb{S} crosses \mathcal{Y} .

Proof:

If a 1-cut doesn't cross \mathcal{Y} , it is included in $\text{scissor}(\mathbb{S}, \mathcal{Y})$ which is a genus 0 surface. Every loop on such a surface is contractible so separating, a contradiction.

□

By calculating the shortest 1-cut passing through every triangle of \mathcal{Y} , we have:

Theorem 3. We can compute a shortest 1-cut in $O(|\mathcal{Y}| \log(V))$.

The algorithm consists in iteratively calculating the shortest 1-cut, while possible. This may not lead to the minimum system of 1-cuts.

2.2 Shortest homotopic cycles on a surface

This second algorithm uses *n-pants* (a genus-0 surface with n boundaries) [example].




Figure 3: A 3-pant

Theorem 4. *In a 3-pant, we can compute the shortest 1-cut **homotopic** to a boundary of the pant in $O(V \log(V))$.*

Proof:

To establish this result, we need (see proof in: [5])

Lemma 4. In a cylinder, we can compute the shortest 1-cut homotopic to a boundary of the cylinder in $O(V \log(V))$

We can use this lemma on a 3-pant by cutting along a shortest path from the two remaining boundaries.

□

In [3], one can find the following result:

Lemma 5. We can compute a first 3-pants decomposition \mathbb{P}_0 (a set of disjoint 1-cuts cutting \mathbb{S} into 3-pants) in $O(gV)$.

Then we can give the pant optimization algorithm: compute a pants decomposition and optimize locally every boundary of a pant, until no improvement is possible.

[3] gives an upper bound for the number of iterations of this algorithm: the (difficult) proof uses rewriting on *crossing words*, words encoding crosses between curves.

In particular, this algorithm is polynomial in its input.

[3] also proves that this algorithm is optimal, in this sense: each cut in the resulting pants decomposition \mathbb{P} is the shortest cut homotopic to the initial corresponding cut in \mathbb{P}_0 . However this algorithm doesn't solve 1-MIN-CUTTING: the homotopy classes of \mathbb{P} are not, *a priori*, equal to the homotopy classes of a minimum system of 1-cuts (it is worth reminding that homotopy classes can't change while optimizing, with this algorithm).

2.3 First attempts to a volume cutting

A first idea is to put volumical information on the surface.

Definition 8. The medial axis (or skeleton) of \mathbb{S} is the set of all points (different from \mathbb{S}) having more than one closest point on the object's boundary, i.e the centers of the open balls tangent to this boundary in at least two points.

Definition 9. The local feature size is the distance to the medial axis.

The method consist in Creating a non-Euclidean distance on the mesh describe by this local feature size: the length of an edge is defined by the local feature size of the adjacent vertices, and by the Euclidean length of the edge, as a product. It corresponds to an estimation of the area that will be introduced if we are correcting an α -junction using this edge.




Figure 4: Local feature size

Another idea, inspired by the second above algorithm, is to decompose a volume into volumic 3-pants and to optimize them. It is the main idea of the approach described next.

3 General algorithm

3.1 Overview

The generic algorithm is the following:

- Compute a first volumic pants decomposition.
- Optimize (i.e change the location while decreasing the weight) every boundary of the cuts while it is possible.
- Among the resulting pants decomposition, select a system of 2-cuts.

Two algorithms result from this pattern, which I call *multi pants (MPA)* and *single pant (SPA) algorithms*.

3.2 First decomposition

To compute a first cutting, I already described the cut locus method: given a set of points A , the *A-cut locus* is the set of all points having more than one shortest path from A .

In the SPA, we simply choose an arbitrary point p , compute the p -cut locus \mathcal{V} and take as an initial decomposition the n -pant which boundaries are the g connected components of \mathcal{V} .

However, one may try to act more *locally* and to use *topological information* to guide the choice of the cuts (i.e to have an initial set of cuts close to the optimal) by decomposing the volume into many pants.

For that purpose, the MPA first compute the skeleton (or medial axis) \mathcal{K} of the surface, using for example an *erosion* process, see [8] for a comprehensive survey of skeleton computation.

Let B be the set of all intersection points of \mathcal{K} : the MPA basically compute the B -cut locus and use its connected components as a pants decomposition.



Figure 5: Left: Skeleton. Right: Cuts associated to this skeleton.

All the resulting pants need not have the same number of boundaries, but assuming that it is the case, we can know precisely the number of such pants:

Theorem 5. *Let p be the number of pants and c the number of cycles of any n -pant decomposition of S .*

Then $p = 2(g-1)$ and $c = 3(g-1)$.

Proof:

For one pant P , according to Euler's formula: $V_P - E_P + F_P = 2 - b = -1$.

Since every couple of pants share a number of vertices equals to the number of edges shared, the number of pants is: $\sum_{P \text{ pant}} (-V_P + E_P - F_P) = V - E + F = 2 - 2g$.

If we duplicate every pant we find $2c = 3g$, since every cut has exactly two adjacent pants.

□

However, we can have intersections in this cut locus, and we are not sure to get 2-cuts [figure].

3.3 Pant(s) optimization

I will first discuss how to optimize a single 2-cut, the next subsection will deal with the problem of selecting a cut to optimize.

We need to distinguish one kind of 2-cut:

Definition 10. A 2-cut \mathcal{C} is an auto 2-cut if the two pants adjacent to \mathcal{C} are actually the same pant.

Since the SPA deals with only one pant, every 2-cuts are auto, and I first describe the optimization of an auto 2-cut \mathcal{C} , associated to a pant \mathbb{P} .

3.3.1 Optimization of an auto 2-cut \mathcal{C}

We first build the network N from G^* the following way:

- Add two vertices s and t to G^* .
- Link s to all the vertices adjacent to \mathcal{C} on one side and the vertices adjacent to \mathcal{C} on the other side to t .

- For each boundary \mathcal{C}' of \mathbb{P} delete \mathcal{C}' , and if $\mathcal{C}' \neq \mathcal{C}$, delete the edges with an end point adjacent to \mathcal{C}' .

The second part of this last point is needed in order to avoid intersections between 2-cuts.

Then, we compute a max flow on N (using Ford Fulkerson or Push-Relabel algorithm, for example) and we use the min cut - max flow theorem (the reader not familiar with flows on network may first read [9]) to deduce a min s - t cut (i.e, a min cut obtained by computing a max flow on \mathcal{C}) \mathcal{C}_{min} . However, \mathcal{C}_{min} is not necessarily a 2-cut, for example $\partial\mathcal{C}_{min}$ is not necessarily connected, but in most applications we hope that \mathcal{C}_{min} is indeed a 2-cut, and from now on, **we assume that every min cut computed this way is a 2-cut**.

It is easy to see that, throughout the SPA, the pant structure is kept ($(\mathcal{V} - \mathcal{C}) \cup \mathcal{C}_{min}$ is a system of 2-cuts).

Moreover, the result of the SPA is likely to be, in practice, independant from the initial set of cuts.

However, it is possible that the best cut \mathcal{C}_{min}' intersects \mathcal{C} , and a max flow will not take it into account.

To avoid this problem, I suggested to take a \mathcal{C}_{min} s-t cut as a new cut, which correctness is validated by the theorem:

Theorem 6. \mathcal{C}_{min}' doesn't intersect \mathcal{C}_{min} .

Proof:

□

3.3.2 Optimization of a 2-cut \mathcal{C} (not auto)




Figure 6: Optimization of a cut

Now we assume that \mathcal{C} is adjacent to two different pants \mathbb{P}_1 and \mathbb{P}_2 , let $\mathcal{C}_{i,j}$ be the boundaries of \mathbb{P}_i different from \mathcal{C} .

We need to be careful since we want to avoid non connected min cut [fig].

We build N from G^* the following way:

- Find a shortest tree \mathcal{T}_1 (resp. \mathcal{T}_2) connecting $\bigcup \mathcal{C}_{1,j}$ (resp. $\bigcup \mathcal{C}_{2,j}$).
- Link s to $\mathcal{C}_{1,j}$ and \mathcal{T}_1 and link $\mathcal{C}_{2,j}$ and \mathcal{T}_2 to t .

This way, every s - t cut is homotopic to \mathcal{C} : **we don't change homotopy classes**.

3.4 Processing choice

Since every optimization of a 2-cut decreases the weight of the 2-cut system \mathcal{Y} , and as there are only finitely many possible total weights (for a discrete volume), our algorithms are *terminating*.

However, the final weight of \mathcal{Y} and the complexity these algorithms depends crucially on the order we choose cuts to optimize. [fig]

Several possible choice were considered:

- Randomly.
- Maximum weighted 2-cut first.
- (All-Min) I also suggested, especially in the SPA, to optimize the 2-cut \mathcal{C} which lead to the minimum possible \mathcal{C}_{min} , and we never optimize \mathcal{C}_{min} .
- Optimize the 2-cut \mathcal{C} maximizing $\text{weight}(\mathcal{C}) - \text{weight}(\mathcal{C}_{min})$.

Bounding the number of iteration required for the 2 first points is likely to be very difficult (personal communication with É. Colin de Verdière).

However the All-Min choice requires $\Omega(g^2)$ max flow, with the SPA: at the beginning there are g 2-cuts to optimize, so we compute g max flow, then we fix one 2-cut so we compute $g - 1$ max flow, and so on.

And indeed, $g + g - 1 + \dots + 1 = \frac{g(g+1)}{2} = \Omega(g^2)$.

Moreover, the validity of the All-Min choice is validated by:

Theorem 7. *Let $\mathcal{C}_1, \dots, \mathcal{C}_k$ be the first k 2-cuts chosen by All-Min \mathcal{C}_k is the minimum [non separating of the decoupage]2-Cut among all cuts not intersecting*

Proof:

Let ... Montrer que il existe cut tel que max flow le donne.

□

3.5 Selection of a cutting system

Proof that we get a topological ball.

4 Improvements

4.1 Neighbourhood

Remark 2. In this case, according to *Menger's theorem* (see [10], corollary 3.3.4), the number of edge disjoint s-t paths is equal to the weight of a min s-t cut.

The most time consuming part of both SPA and MPA being the computation of a max flow, one may try to improve its time complexity, and indeed, we can act more locally (for the importance of local search and neighbourhood, see [11]).

Definition 11. If v is a vertex and p a path, $d(v, p)$ (and $d(p, v)$) is the minimal length of a shortest path between v and a vertex of p .

Given two paths p and p' , $d(p, p') = \max_{v \in p} d(v, p')$.

The p -neighbourhood \mathcal{N}_p is the set of all paths p' such that $d(p, p') \leq 1$.
The strict p -neighbourhood $\partial\mathcal{N}_p$ is the set of all paths p' such that $d(p, p') = 1$.

If $\mathcal{P} = \{p_1, \dots, p_n\}$ is a set of paths, $\mathcal{N}_{\mathcal{P}} = \Delta \mathcal{N}_{p_i} (= (\bigcup \mathcal{N}_{p_i}) - \bigcup p_i)$, and we define similarly $\partial\mathcal{N}_{\mathcal{P}}$.

Remark 3. d is a distance: $d(p, p') = d(p', p)$, $d(p, p') = 0 \iff p = p'$ and $d(p, p'') \leq d(p, p') + d(p', p'')$.

Theorem 8. *Assume we are computing a max flow in \mathbb{V} , using Ford Fulkerson algorithm and that we already found n augmenting paths $\mathcal{P} = \{p_1, \dots, p_n\}$ with BFS.*

Let R be the residual network, and f the current flow.

If there is no path from s to t in $R \cap \mathcal{N}_{\mathcal{P}}$ then f is a max flow.

Proof:

□

Remark 4. If we allow capacities to be different, we have the following weaker statement: if there is no path from s to t in $R \cap \mathcal{N}_{\mathcal{P}}$ then f is a max flow.

Then, to compute a max flow, we can use the following algorithm:

```

 $\mathcal{P} \leftarrow$  shortest path from  $s$  to  $t$  in  $R$  (using BFS) ;
while There is an augmenting path  $p$  from  $s$  to  $t$  in  $R \cap \partial\mathcal{N}_{\mathcal{P}}$  do
   $\mathcal{P} \leftarrow \mathcal{P} \cup p$  ;

```

4.2 Processing order and intersections

4.3 Complexity

5 Implementation

5.1 Overview

Initially, it was planned that I develop the optimization part of the MPA, and that the rest of the team will develop the initial pants decomposition, all in C++.

After two weeks working on the core algorithm, I decided to use tetrahedra, although more complicated to implement than voxels, to be able to work on surfaces, on Blender for example and use TetGen to tetrahedralize them.

I also used TetMeshLib, a library to manage tetrahedral meshes that Marco Attene developed.

To both automatize the different tools of the algorithm and visualize more easily, I wrote a script in Blender Python API.

At the end of the internship, I worked together with Jean Marie Favreau, using GitHub.

Globally, I wrote about 3000 lines in C++ and 400 in Python and the documentation, built with doxygen, is available at [url].

5.2 SGL

SGL relies on *generic programming*: an ADT (Abstract Data Type) defines the required methods and semantic for a class and is used through template parameter.

For example, the class BFS (Breadth First Search) has a Graph template parameter, which must provide, among other, a class named iterator giving a way to iterate over all edges adjacent to a given vertex.

OC3D was designed in a similar way.

5.3 Data structures




Figure 7: Hierarchy of edge classes (an arrow means "inherits from")

Since we often need to iterate over all edges adjacent to a vertex, every graph is implemented as a adjacency list graph (`Graph_List<Edge>`).

I use mainly three graphs: the pants graph, the dual graph G^* and the dual-adjacency graph $\overline{G^*}$.

The pants graph is directed, while G^* and $\overline{G^*}$ are undirected, this is needed by the way the residual graph is used by the implementation of Ford Fulkerson algorithm: we access the residual capacity of an edge of G^* (of type `Edge_Dual`) providing the direction of the edge (i.e, the end extremity), it uses less memory and time since we manipulate only one edge pointer (this edge pointer is stored in the adjacency list of both extremities).

Since the network we use for Ford Fulkerson algorithm is not oriented, every edge of G^* has a reverse (in the opposite direction) and every edge has a pointer to its reverse.

An edge (i.e, a 2-Cut) of the pants graph is of type `Edge_Cut`: its extremities are the pants adjacent to the corresponding 2-Cut and it stores the list of the pointers to `Edge_Dual` elements included in the 2-Cut.

Every `Edge_Dual` in a `Edge_Cut` has the same orientation, i.e their end extremities are in the same pant.

This is needed to know to which vertices we must link the source (or sink) before running Ford Fulkerson algorithm, or to delete the edges out of the pant we optimize (subsection 5.5 explains how to orientate a cut).

To avoid time waste, every 2-Cut has a reverse and a pointer to it.

5.4 Core algorithm




Figure 8: Hierarchy of IO classes.

To improve genericity, the algorithm first constructs the dual graph (which depends on the representation used - voxels or tetrahedra) and then works only on the dual graph.

For that purpose, an object is used to deal with representation-dependant operations, `IO_Base` implements the required methods of such an object but the dual graph must be set manually, `IO_Tet` defines how to make the dual graph from tetrahedra and `IO_Voxels` from voxels image.

`IO_Tet_Adj` also defines how to build the adjacency-list graph.

After the dual graph is made, we can search a max flow on it with a class derived from `Max_Flow`, and `Cut_vertices` is a class giving a min cut from a max flow, searching for the set S of every vertices reachable from the source s and setting as the cut the set of edges connecting S and S^c .




Figure 9: Simplified hierarchy of max flow classes.

5.5 Orientation

When we get a cut, selecting it with Blender for example, we must orientate it. After discussing it with Jean Marie Favreau and Vincent Barra, the best solution may be the following:

- Take an element of the cut (a triangle or a square, say a triangle abc), give an orientation to it, which can be viewed as a permutation σ of its vertices.
- For each triangle uvw adjacent to it with an edge uv , if $\sigma(u) = v$, defines the orientation on t according to the permutation vuw .
- Repeat the second point replacing abc by all new triangles with an orientation, until every triangle has an orientation.
- For each edge uv in the cut, associated to a triangle with orientation (abc) , compute $\det(a,b,c)$: if it is negative, replace uv by vu .

This algorithm is clearly linear (assuming that the number of vertices of an element of a cut is fixed).




Figure 10: An orientation.

5.6 Python script

The script (which can be found together with a video describing its working) I wrote for Blender uses, with pipeline, OC3D_debug, a command-line debugger which can, among other, load a mesh, make the dual graph, optimize cuts step by step and output various information.

To avoid freezing Blender while optimizing a cut I used several threads and I had to synchronize the threads, OC3D_debug and TetGen.




Figure 11: Blender Python API interface.

5.7 Running time comparison

To compare the algorithms of max flow, I tested them on a slightly modified torus, with various number of vertices.

The examples can be found at : [\[url\]](#).

The timing runs were performed on a Dell inspiron 1520 with Intel Core 2 Duo T7300 Processor (2.0GHz, 4MB L2 cache) and 2048 Mo memory.

Compilation was done with Visual C++ 2008 with all optimization flags set for maximum speed.

The times require by the continue neighborhood algorithm is of same order of the time required by TetGen (about 80 seconds for the torus which dual has a complexity 1750000).

The advantage of a neighborhood increases significantly as the number of vertices of G^* increases.

Complexity (edges + vertices) of the dual graph	Classic	Neighborhood	Continue
200K	111s	61s	29s
900K	1147s	236s	101s
1750K	2975s	318s	158s

Table 1: Running time in seconds of the optimization of one 2-Cut using several algorithms of max flow (K stands for 1000).

I also compared, on the dual graph with complexity 900000, the time required for a BFS at different moments of the algorithms.

I split the execution of the algorithms in different steps defined by the augmentation of the neighborhood used by the continue neighborhood algorithm (thus, the continue neighborhood algorithm found 1 path before augmentating, then 21, 129 and so on).

Then I compared the number of edges of the neighborhood in which the paths were found, the total time during the step, and the average (shortened Avg) time of a BFS.

The advantage of a neighborhood decreases slightly during the algorithm, since the neighborhood increases as the number of edges visited by a BFS on the whole graph decreases (some edges have then zero capacity and disappear from the residual graph), so the algorithm is likely to be more efficient if the min cut is small (this is especially true for voxels with capacity one, since the number of iterations is directly the number of iterations of the algorithm).

	Continue neighborhood			Neighborhood			Classic	
Paths found	Edges	Total time	Avg time	Edges	Total time	Avg time	Total time	Avg time
1	0	250	250	0	250	250	250	250
21	2452	41	1,95	2452	49	2,33	4917	234,14
129	4240	283	2,19	5084	446	3,46	30140	233,64
363	9428	1295	3,57	11282	2500	6,89	86221	237,52
741	21822	5234	7,06	25930	10028	13,53	165232	222,99
1735	44944	22627	13,04	50886	46345	26,71	364113	209,86
2798	87172	66810	23,88	101262	144944	51,8	559327	199,9
249	161664	10568	42,44	184092	16781	67,39	48611	195,22
0	173368	22	22	189654	33	33	94	94

Table 2: Detailed time comparison

6 Appendix

6.1 Mathematical background

Definition 12. is a continuous map $\gamma : [0, 1] \longrightarrow \mathbb{S}$ such that $\gamma(0) = \gamma(1)$

Warning: non separating neq contractible




Figure 12: Average time of a BFS with respect to the step in the CNA.

6.2 Applications

6.3 SGL

References

- [1] J.-M. Favreau. Outils pour le pavage de surfaces. PhD thesis, Blaise Pascal university, 2009.
- [2] É. Colin de Verdière. Raccourcissement de courbes et décomposition de surfaces. PhD thesis, Paris 7 university, 2003.
- [3] É. Colin de Verdière and F. Lazarus. Optimal pants decompositions and shortest homotopic cycles on an orientable surface. Proceedings of the 11th Symposium on Graph Drawing , 2003.
- [4] J.-M. Favreau, V. Barra. Conformal Flattening of the Cortical Surface using a Topological and Geometrical Cutting. SEECCM, South-East European Conference on Computational Mechanics, Island of Rhodes, Greece, 2009.
- [5] É. Colin de Verdière. Algorithms for graphs on surfaces, M.P.R.I. course Master 2. 2009-2010.
- [6] Cortical surface mapping using topology correction, partial flattening and 3D shape context-based non-rigid registration for use in quantifying atrophy in Alzheimer’s Disease. Preprint submitted to International Journal of Biomedical Imaging, March 2010.
- [7] J. Erickson, S. Har-Peledz. Optimally Cutting a Surface into a Disk. Discrete and Computational Geometry, July 2, 2002.
- [8] N. D. Cornea, D. Silver, P. Min. Curve skeleton properties, applications and algorithms. IEEE Transactions on Visualization and Computer Graphics 13 (2007), pp 530-548.
- [9] T. Cormen, C. Leiserson, R. Rivest. Introduction to algorithms. MIT Press, 1990.
- [10] R. Diestel. Graph Theory, volume 173 of Graduate Texts in Mathematics. Springer-Verlag, 2000.

- [11] C. H. Papadimitriou, K. Steiglitz. Combinatorial optimization: Algorithms and Complexity. Dover Publications, 1998.

Proof:

□